

Ensemble Methods

November 20, 2018

1 Ensemble methods demo

This notebook shows how ensemble methods can improve results in a practical manner.

First, we are training two models - one decision tree and one logistic regressor. We train these on the classic [wine quality dataset](#) by Cortez et.al, which is a regression problem to predict the wine quality based on some features such as fixed acidity, citric acid, chlorides and alcohol. Then, we combine the results and look at the mean squared error between the models predicted output and the test set, compared to each models output separately.

```
In [124]: import pandas as pd
          from sklearn import tree
          from keras.layers import Input, Dense
          from keras.models import Model
          from sklearn.model_selection import train_test_split
          from sklearn.linear_model import LogisticRegression
          import numpy as np

          from matplotlib import cm
          import matplotlib.pyplot as plt

          plt.style.use('ggplot')
          %matplotlib notebook
          from mpl_toolkits.mplot3d import Axes3D
```

We can see that the dataset contains a lot of interesting features, and that there are a total of 4898 entries in the dataset.

```
In [149]: data = pd.read_csv('winequality-white.csv', sep=";")
          data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4898 entries, 0 to 4897
Data columns (total 12 columns):
fixed acidity          4898 non-null float64
volatile acidity       4898 non-null float64
citric acid            4898 non-null float64
residual sugar         4898 non-null float64
chlorides              4898 non-null float64
```

```

free sulfur dioxide      4898 non-null float64
total sulfur dioxide     4898 non-null float64
density                  4898 non-null float64
pH                       4898 non-null float64
sulphates                4898 non-null float64
alcohol                  4898 non-null float64
quality                  4898 non-null int64
dtypes: float64(11), int64(1)
memory usage: 459.3 KB

```

```
In [74]: # Load X and Y data and split into training and test set
```

```

Y = data['quality']
X = data.drop(['quality'], axis=1)
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.33, random_state=42)
print("Shape of x train:", X_train.shape)
X.head()

```

```
Shape of x train: (3281, 11)
```

```

Out[74]:
fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
0             7.0                0.27        0.36              20.7        0.045
1             6.3                0.30        0.34              1.6        0.049
2             8.1                0.28        0.40              6.9        0.050
3             7.2                0.23        0.32              8.5        0.058
4             7.2                0.23        0.32              8.5        0.058

free sulfur dioxide  total sulfur dioxide  density  pH  sulphates \
0                 45.0                 170.0  1.0010  3.00      0.45
1                 14.0                 132.0  0.9940  3.30      0.49
2                 30.0                 97.0   0.9951  3.26      0.44
3                 47.0                 186.0  0.9956  3.19      0.40
4                 47.0                 186.0  0.9956  3.19      0.40

alcohol
0      8.8
1      9.5
2     10.1
3      9.9
4      9.9

```

```
In [151]: tree_model = tree.DecisionTreeClassifier()
log_reg = LogisticRegression(random_state=0, solver='lbfgs', multi_class='multinomial')
```

```
In [152]: # Fit both models
```

```

tree_model.fit(X_train, Y_train)
log_reg.fit(X_train, Y_train)

```

```
/home/stud2173/anaconda3/envs/ml/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:
    "of iterations.", ConvergenceWarning)
```

```
Out[152]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=1000, multi_class='multinomial',
    n_jobs=None, penalty='l2', random_state=0, solver='lbfgs',
    tol=0.0001, verbose=0, warm_start=False)
```

```
In [153]: # Get both models to make predictions on the test set
y_pred_1 = tree_model.predict(X_test)
y_pred_2 = log_reg.predict(X_test)
```

```
In [154]: def MSE(y_pred, y):
    N = len(y_pred)
    return 1/N * sum((y_pred - y)**2)

    tree_loss = MSE(y_pred_1, Y_test)
    log_reg_loss = MSE(y_pred_2, Y_test)
    print("Loss for decision tree model: ", tree_loss)
    print("Loss for logistic regression model: ", log_reg_loss)
```

```
Loss for decision tree model:  0.7260358688930117
Loss for logistic regression model:  0.6474953617810761
```

```
In [155]: # By taking the average of both models, we get a better prediction
averaged_pred = 0.5*y_pred_1 + 0.5* y_pred_2
avg_loss = MSE(averaged_pred, Y_test)
print("Averaged prediction loss: ", avg_loss)
```

```
Averaged prediction loss:  0.509121830550402
```

2 Searching for best ensemble weighing

Instead of taking the average of both models, we can search the weight space for the best weighing between the two models. Here, we are searching with an increment of 0.05.

Since the weighed total is 1, we only have to plot one of the weights - w_1 . The other weight is $1-w_1$. We observe that the weights favor the logistic regressor (which is not surprising given that it has a bit higher accuracy).

```
In [156]: WEIGHTS = np.linspace(0, 1, 21)
LOSS_VALS = []
minima = 10000
best_weights = (0, 0)
for w1 in WEIGHTS:
    w2 = 1-w1
```

```
weighted = w1 * y_pred_1 + w2*y_pred_2
```

```
loss = MSE(weighted, Y_test)
```

```
if loss < minima:
```

```
    best_weights = w1, w2
```

```
    minima = loss
```

```
LOSS_VALS.append(loss)
```

```
print("Found best weights to be {} with loss {}".format(best_weights, minima))
```

Found best weights to be (0.45, 0.55) with loss 0.5069712430426709

```
In [148]: plt.plot(WEIGHTS, LOSS_VALS, '-')
```

```
plt.title("Weighted ensemble loss")
```

```
plt.xlabel("w1")
```

```
plt.ylabel("MSE")
```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

Out[148]: Text(0,0.5,'MSE')